

Microprocessor IO Subsystems



Chun-Jen Tsai
NYCU
12/09/2022

Interface btw. μ P and Devices

- ❑ There are several methods to design communication interface between processors and devices of a system:
 - Dedicated I/O instructions – *obsolete*
 - Memory-mapped I/O interface
 - Interrupt-driven interface – *mailbox device*
 - Co-processor interface – *ARM does it, with FP Unit*

- ❑ In any cases, the set of wires that connects HW modules together is called a bus

Dedicated I/O Instructions

- ❑ In early CPU ISA designs, device I/O and memory R/W are done using different instructions
- ❑ For example, in Intel x86 ISA:

device I/O

```
IN  BX, 19h
ADD BX, 1
OUT 19h, BX
```

device I/O

v.s.


Memory R/W

```
MOV BX, [19h]
ADD BX, 1
MOV [19h], BX
```

memory I/O


load/store in RISC

- In previous example, I/O address 19h and memory address [19h] points to different space.

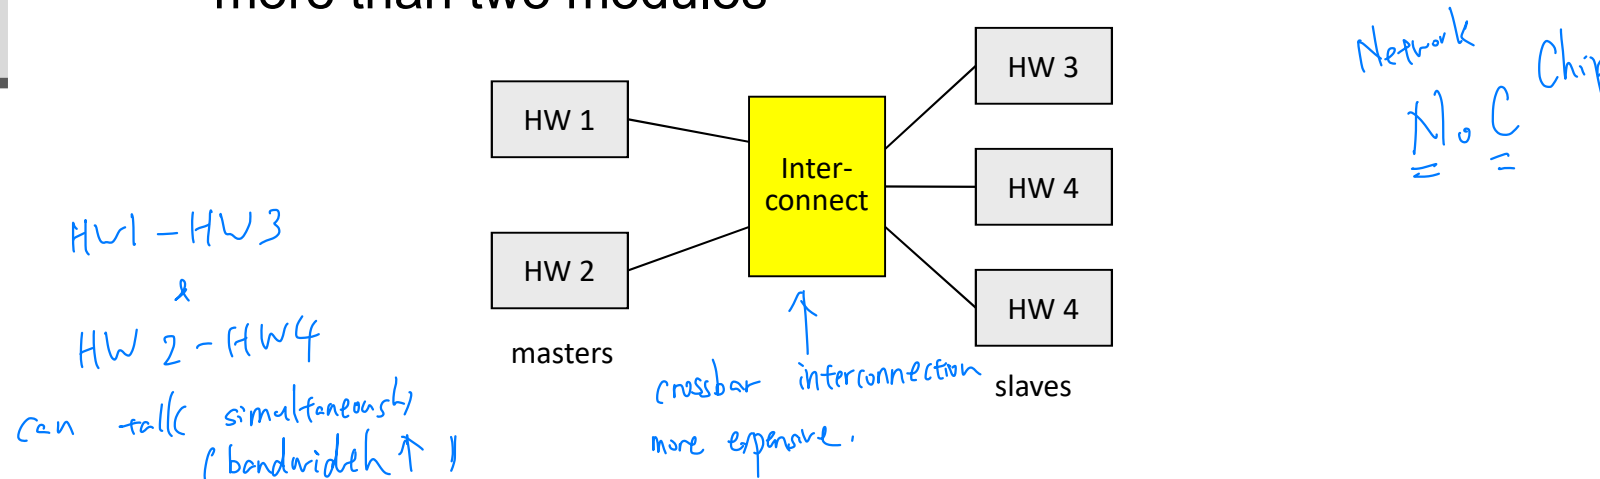


Bus Topology


- ❑ ^{Ethernet} Shared bus – a set of wires is shared among modules
 - Shared bus saves both design and production costs
 - Performance of the system may suffer



- ❑ Point-to-point bus – dedicated links btw. two modules
 - Complex inter-connect module is required when there are more than two modules



Bus Components



❑ Bus master:

- A HW module that is able to initiate read and write requests
- Typical masters: CPU and DMA

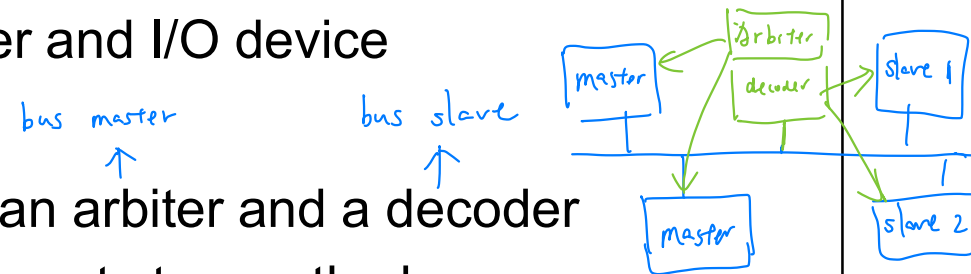
Direct Memory Access (L/S large chunk of memory)

❑ Bus slave:

- Responds to read/write requests from the bus masters
- A slave signals back to the master the results of data transfer
- Typical slaves: memory controller and I/O device

❑ Bus controller:

- A bus controller is composed of an arbiter and a decoder
- Arbiter: determines which master gets to use the bus
- Decoder: determines which slave should reply to a request





Although without a decoder, the architecture can still work, the slave must listen to the bus, waiting for a request. With a decoder, the slave IP can be turned on only when needed, saving power.

Bus Protocols



- ❑ Bus protocol controls communication among modules:
 - Determines who can use the bus at any particular time
 - Governs length, style of communication

- ❑ Data framing vs. control wires
 - A bus architecture with few wires usually uses “data framing” (at data-link layer) to signal transactions
 - Control and data signals shares the same wires
 - Sometimes referred to as “in-band” signaling
 - A bus architecture with many (parallel) wires uses “control wires” (at physical layer) to signal transactions
 - Control and data signals use different wires

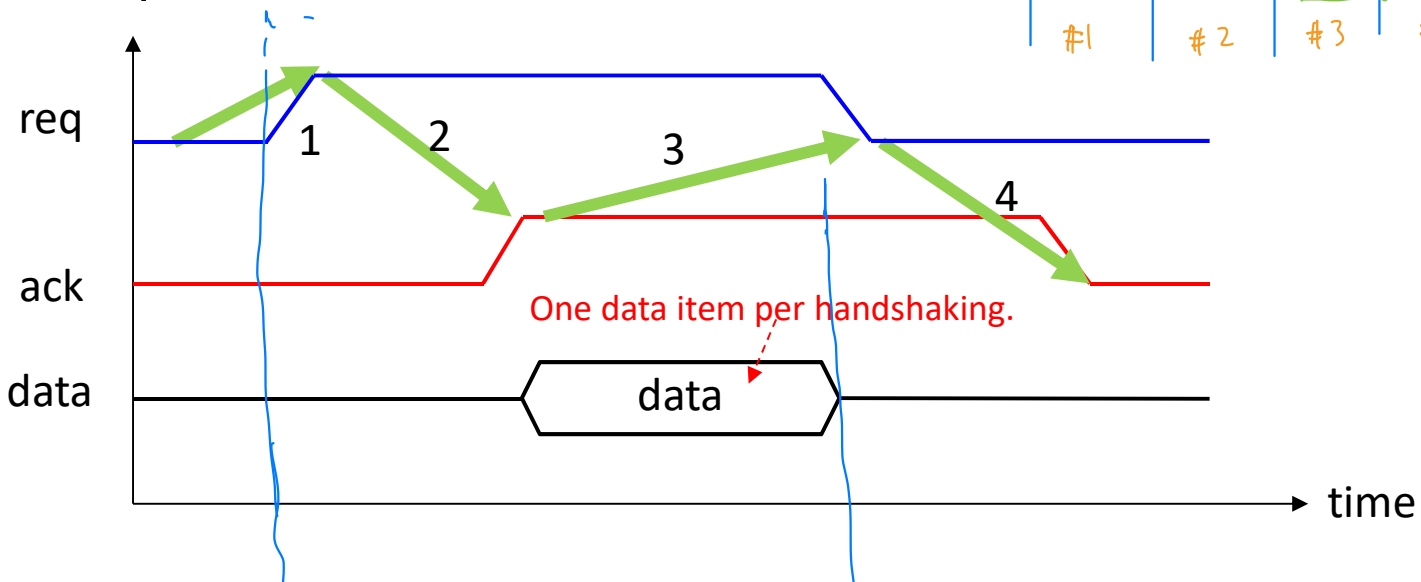


Four-Cycle Bus Protocols

- ❑ Four-cycle handshake is the basis of many protocols
 - Use two control wires: req (request) and ack (acknowledgment)




- Example of transaction:



may be multiple
data lines eg
16 data

req send




Formal BUS Behavior

- ❑ Formal properties of a BUS protocol:
 - Following a reset, everything should be set to idle
 - If a request is stalled (hit wait state), it should be maintained
 - There shall be no responses without prior requests
 - All requests get responses
 - Once a transaction is done, everything should return to idle

A request must maintain because the IP of bus bridges may latch the request.

Interface Using Memory-Mapped I/O

- ❑ Processor interface can be more flexible with the memory mapped I/O (MMIO) model
- ❑ In MMIO model, some “memory” addresses are mapped to the internal control registers of the device
 - Processor write data into these addresses to control the device
 - Processor read data from these addresses to retrieve results from the device



Typical Bus Signals for MMIO

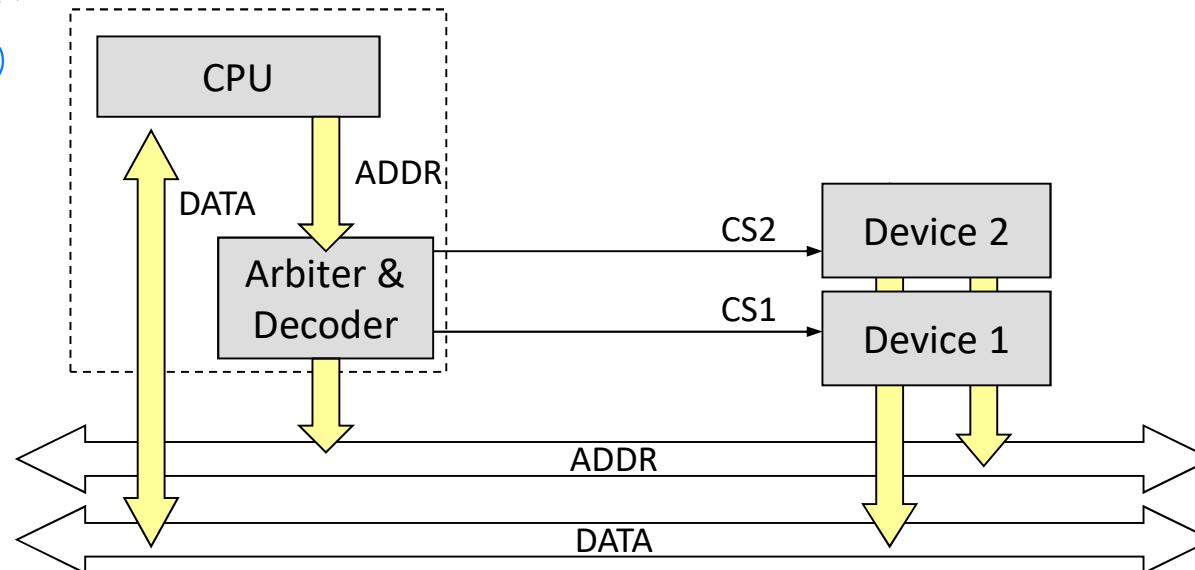
FF ~~XX~~ ID for slave

- ❑ ADDR: multi-bit ID for slave devices
- ❑ RW: signals a read or write request
- ❑ DATA: multi-bit data signal (to or from slave)
- ❑ READY: from slave (ack)
- ❑ CS: device selection signal (for power saving)

Request from master

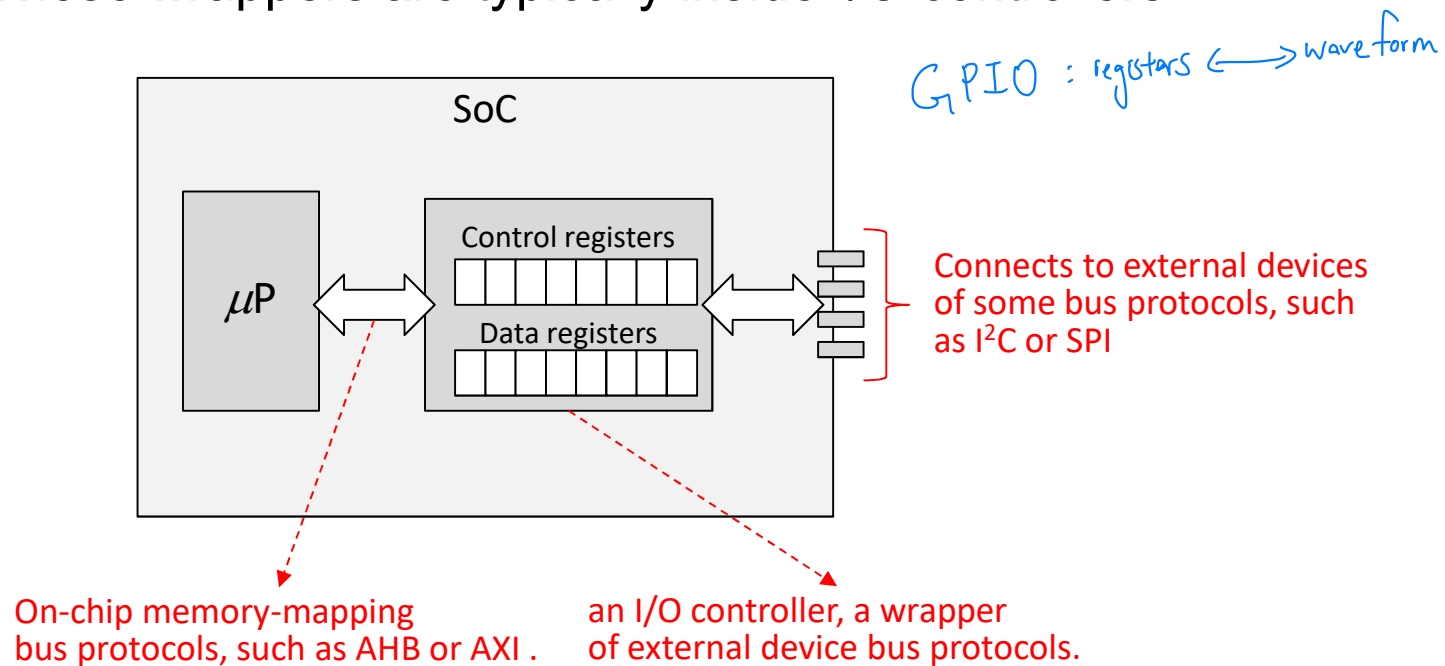
Ack from slave

chip-select-
(en, enable)




MMIO Wrapper for Devices

- ❑ Since CPU uses load/store instructions to talk to various devices, we need a wrapper circuit to help converting read/write operations to bus signals
 - These wrappers are typically inside I/O controllers




μ P Interface Using GPIO Pins

- ❑ Most microcontrollers have some GPIO pins that can be used for interfacing with the device logic
 - GPIO is a simple HW device that maps a register bit value to a input/output wire of the CPU pin



*GPIO could talk to low speed devices
e.g. SD-card, I/O2*

- ❑ The number of GPIO pins of a CPU is usually small
 - Multiplexing technique can be used to issue a request



Talking to ASIC from C via GPIOs

- ❑ Typical GPIO device contains a n -bit register
 - Each register is connected to a I/O pin
 - Physical pins are implemented as “inout” ports

- ❑ C code example:

```
unsigned int *PIO = 0x03FF5008; // GPIO register address


/* write */
*PIO = 0xB0;

/* read */
Data = (*PIO >> 8) & 0xFF;
```

Inter-Integrated Circuit

I²C Bus

Originally designed for TV controllers

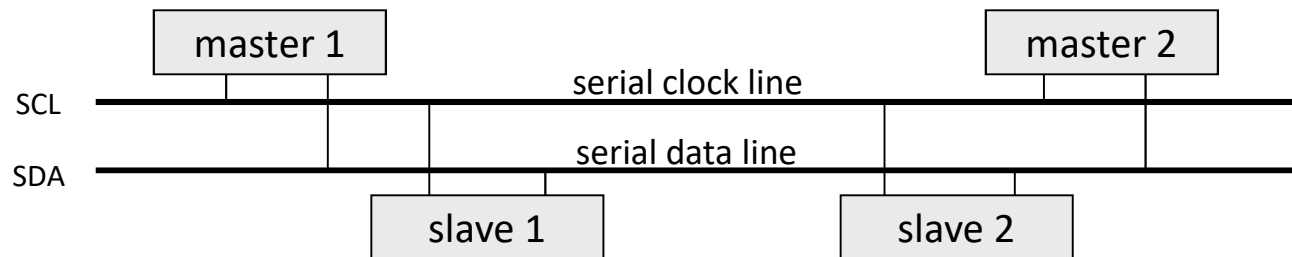


❑ Designed by Phillips in 1982

- Two-wire, synchronous serial link from 100kHz ~ 5MHz
- Shared bus with multi-master support using fixed-arbitration
- Most SoC's come with built-in I²C controllers
 - Boot code stored in an I²C flash, and read into RAM at power-on

❑ I²C signaling

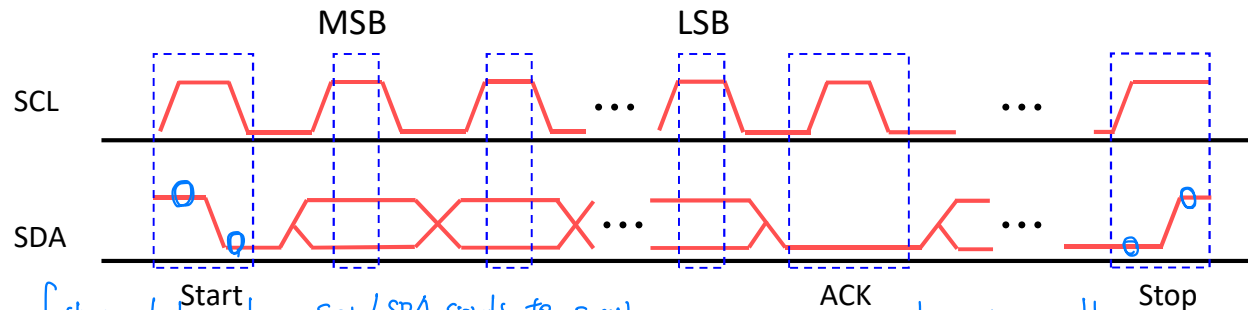
- SDA carries 1 by default, sender pulls it down to send 0
- If sender sends a 1 and gets a 0, others are sending data



Any master who detects a collision shall not keep sending


I²C Signaling

- ❑ I²C transmissions occur in 8-bit bytes
 - Master raises a START pattern
 - Send 8 bits (slave ID & R/W, or data)
 - Target slave responds with an ACK pattern
 - Master sends a STOP pattern when the transaction is done



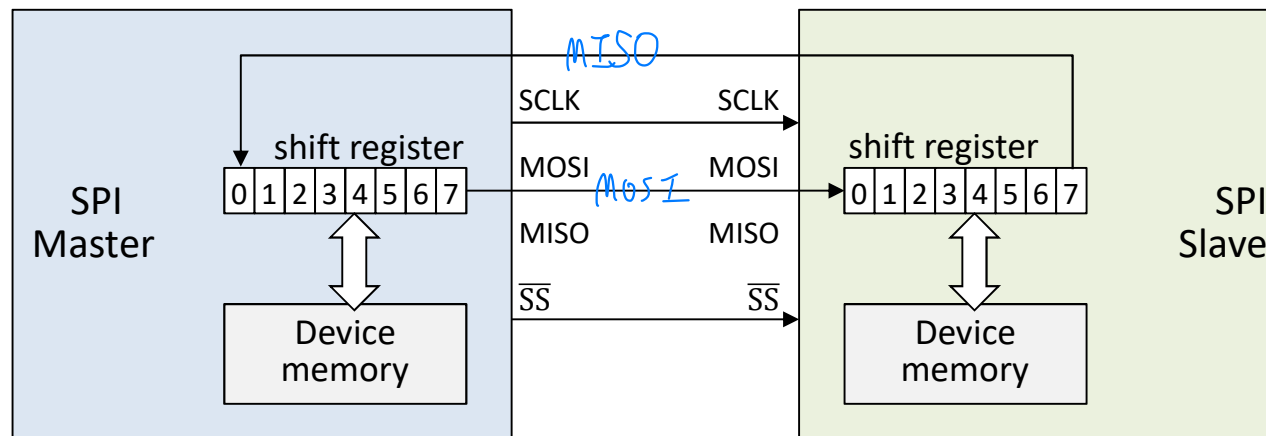
We need a faster clock ^{Start} than SCL/SDA signals to sample points on them, this is possible ^{ACK} ^{Stop}

- ❑ An example of transmission: *because I²C is a low speed protocol.*



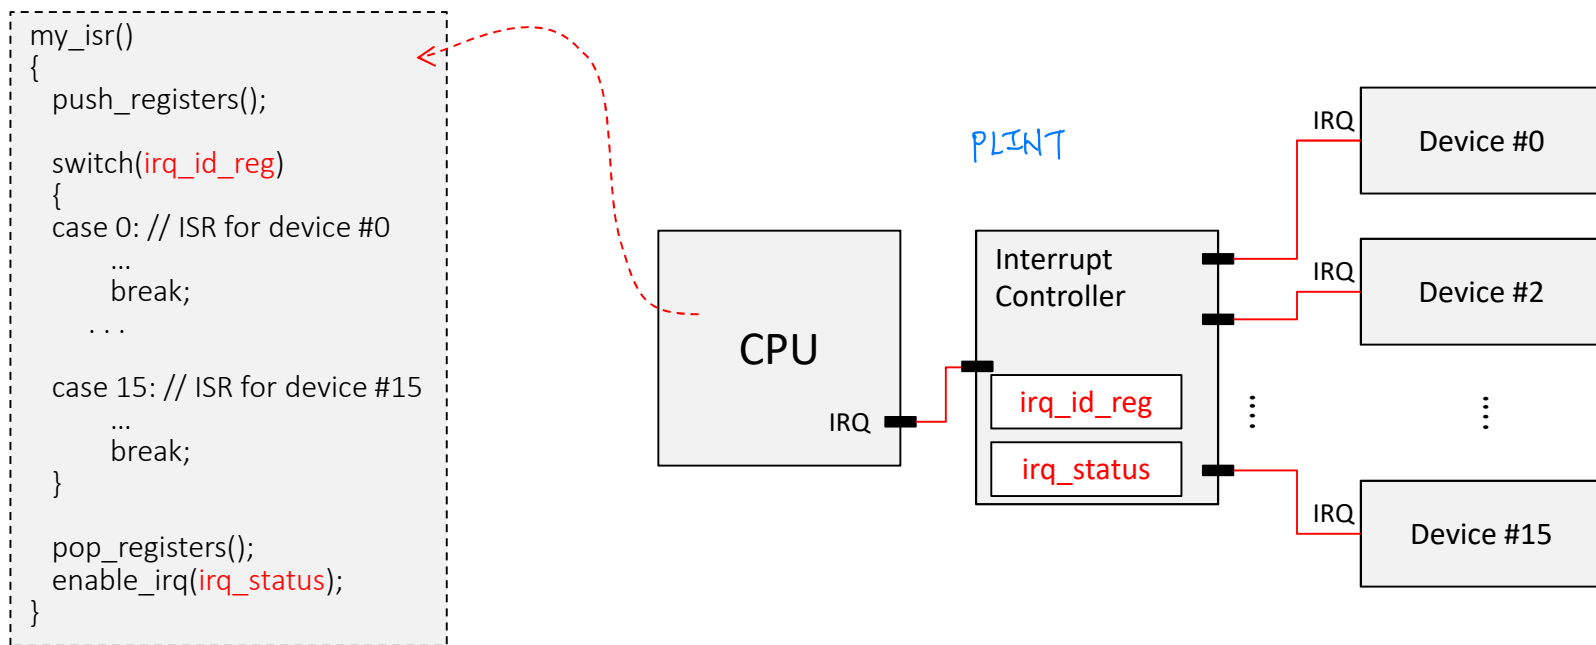
Serial Peripheral Interconnect (SPI)

- ❑ Designed by Motorola in 1980's
 - Short-distance, full-duplex, synchronous serial transmission
 - One master, multiple slaves
 - Often used for SD cards, flash ROMs, etc.
 - Master selects the target slave via \overline{SS} (slave selection)
 - Data sizes are 8, 12, or 16 bits per transaction



Interrupt Signaling from IP to CPU

- ❑ Sometimes, a HW module sends an interrupt to CPU to signal a request or an acknowledgment (ack)
- ❑ Typical system diagram with interrupt-capable devices



Interrupt Service Routine (ISR)


Example: RISC-V Interrupt Interface

- ❑ A RISC-V core has 3 execution modes: M, U, and S
 - Only M mode, the machine mode, is mandatory.
 - Current execution mode is stored in a private register
- ❑ Interrupt info is in the Control & Status Registers (CSR)
 - Each mode has its own set of registers
- ❑ The execution mode can only be changed by a trap, an interrupt, or `mret/sret/uret` return instructions
 - The core stores the interrupted mode, and resume its value upon return

RISC-V Interrupt Flow

- ❑ For example, for M-mode interrupt:
 - Store the interrupted PC in a CSR register `mepc`
 - Set the interrupt cause in a CSR register `mcause`
 - Disable any interrupts by flagging the `mie` bit in `mstatus`
 - Look up the interrupt handler address in `mtvec`:
 - `mtvec` is composed of `{BASE[31:2], MODE[1:0]}`
 - In direct-mode, `BASE` is the ISR address
 - In vector-mode, the `BASE+4×mcause` is the ISR address
 - Transfer control to the ISR

The ISR may jump to a jump-table first, then jump to the correspond ISR



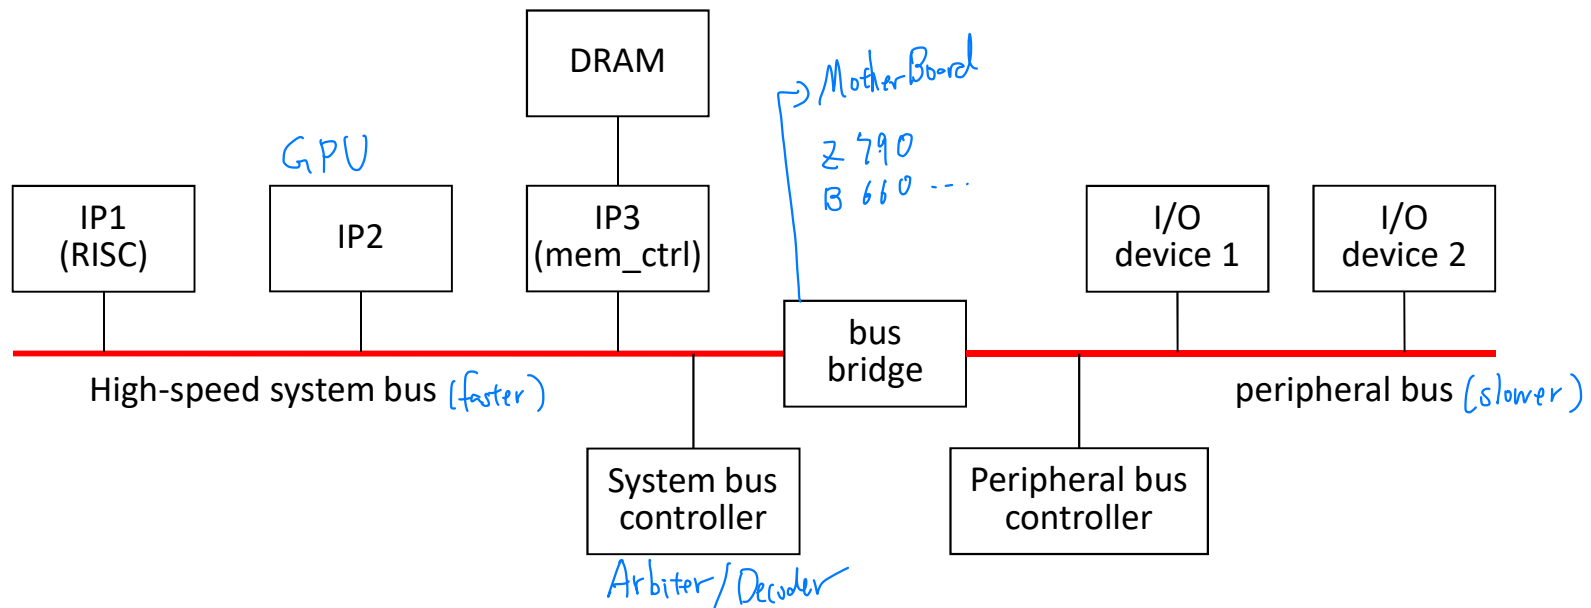
On-Chip Bus Architecture

- ❑ On-chip bus is one of the crucial components for SoC
 - Facilitates platform-based design paradigm
 - Enables reusable IP design *platform based SoC Design.*

- ❑ Popular SoC bus architecture:
 - AMBA – ARM's bus architecture
 - AHB/APB: shared bus topology
 - AXI: point-to-point bus topology
 - CoreConnect – IBM's bus architecture (for PowerPC)
 - Wishbone – the OpenCores bus for open source IP design

Typical SoC Bus Architecture

- ❑ A traditional bus architecture (e.g., ARM's AHB and APB) are designed for the following system model:




- ❑ Today, the buses often become the system bottleneck!

NoC: Put a mini Network on Chip to interconnect IPs.


Use wireless communication inside a chip.

Physical Separation of I/O Busses


- ❑ For high-speed buses, the input and output lines are usually separated into two unidirectional buses
 - The system bus controller (contains the bus arbiter and bus decoder logics) will establish a unidirectional high-speed link from the data source IP to the data destination IP:



Share bus
We need a tri-state buffer



if "en" is turned off, it's like disconnect wire



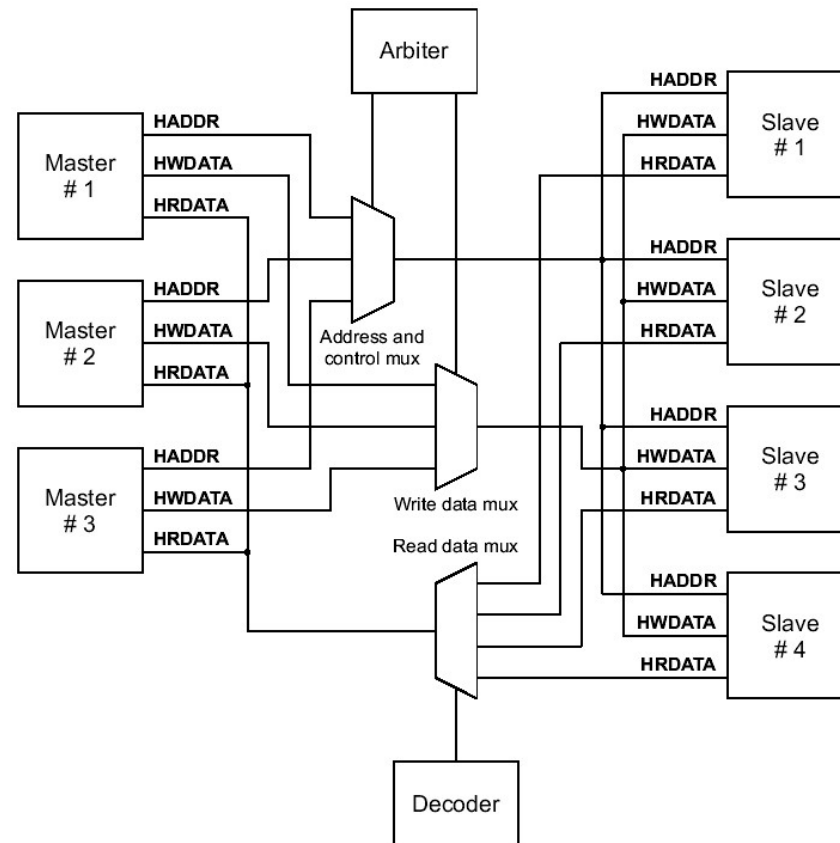
drawback: low clock rate

ARM Bus Evolution

- ❑ The Advanced Microcontroller Bus Architecture (AMBA) is defined by ARM, Ltd
- ❑ Versions of AMBA:
 - AMBA 1.0 – ASB and APB, 1996
 - AMBA 2.0 – **AHB** and APB, 1999 *AHB: High performance unidirectional bus*
 - AMBA 3.0 – AXI 1.0 and ATB (for CoreSight), 2003
 - AMBA 4.0 – **AXI4** (aka, AXI 2.0) and ACE, 2010-2011
 - AMBA 5.0 – CHI (Coherent Hub Interface), 2013
- ❑ ARM buses are free of licenses

AHB Multiplexer Interconnection

- ❑ AHB is a multiplexed bus with a centralized arbiter

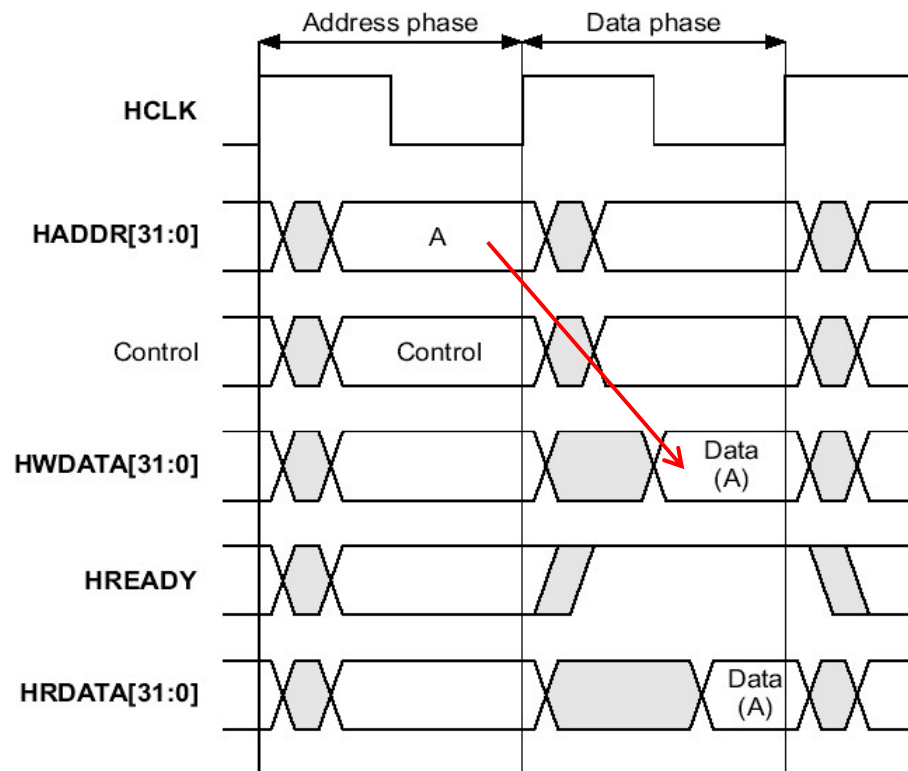


† Figure extracted from: ARM Limited, *AMBA Specification Rev 2.0*, ARM IHI 0011A, May 13, 1999, figure 3-2, page 38.

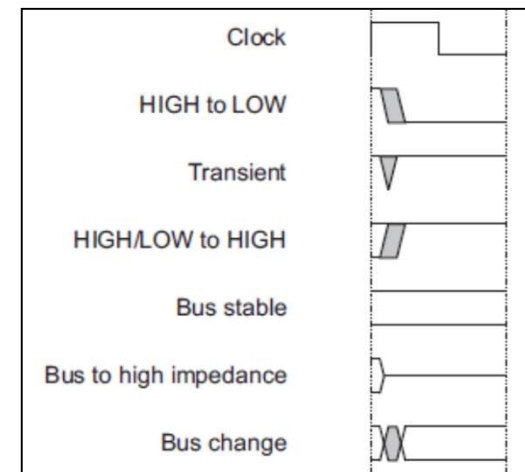
AHB Simple Transfer

(- 単打)
 single-beat transfer

- Each transfer consists of an address/control cycle and one or more data cycles:

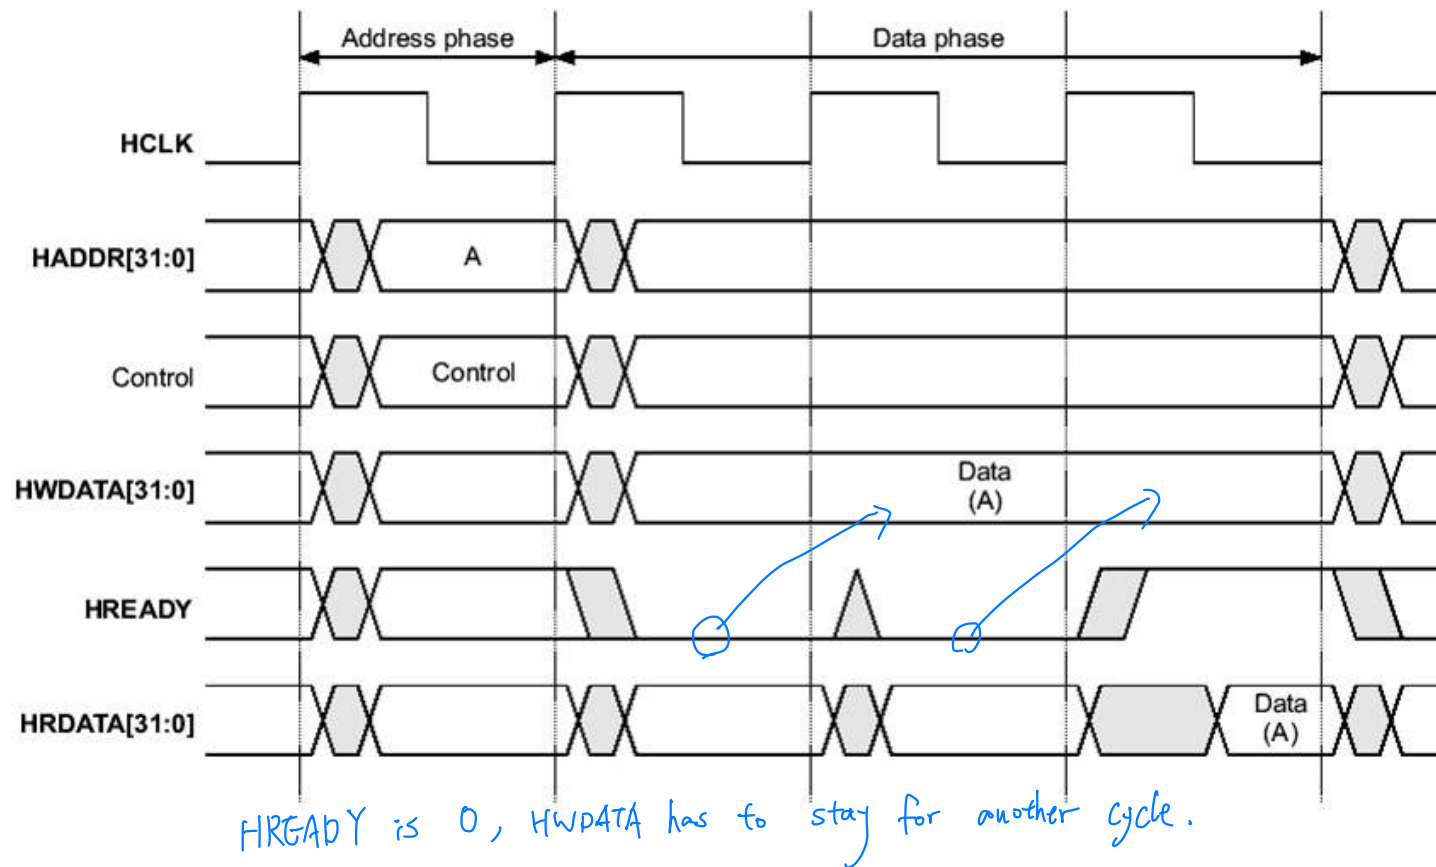


Bus diagram convention:



† Figure extracted from: ARM Limited, *AMBA Specification Rev 2.0*, ARM IHI 0011A, May 13, 1999, figure 3-3, page 40.

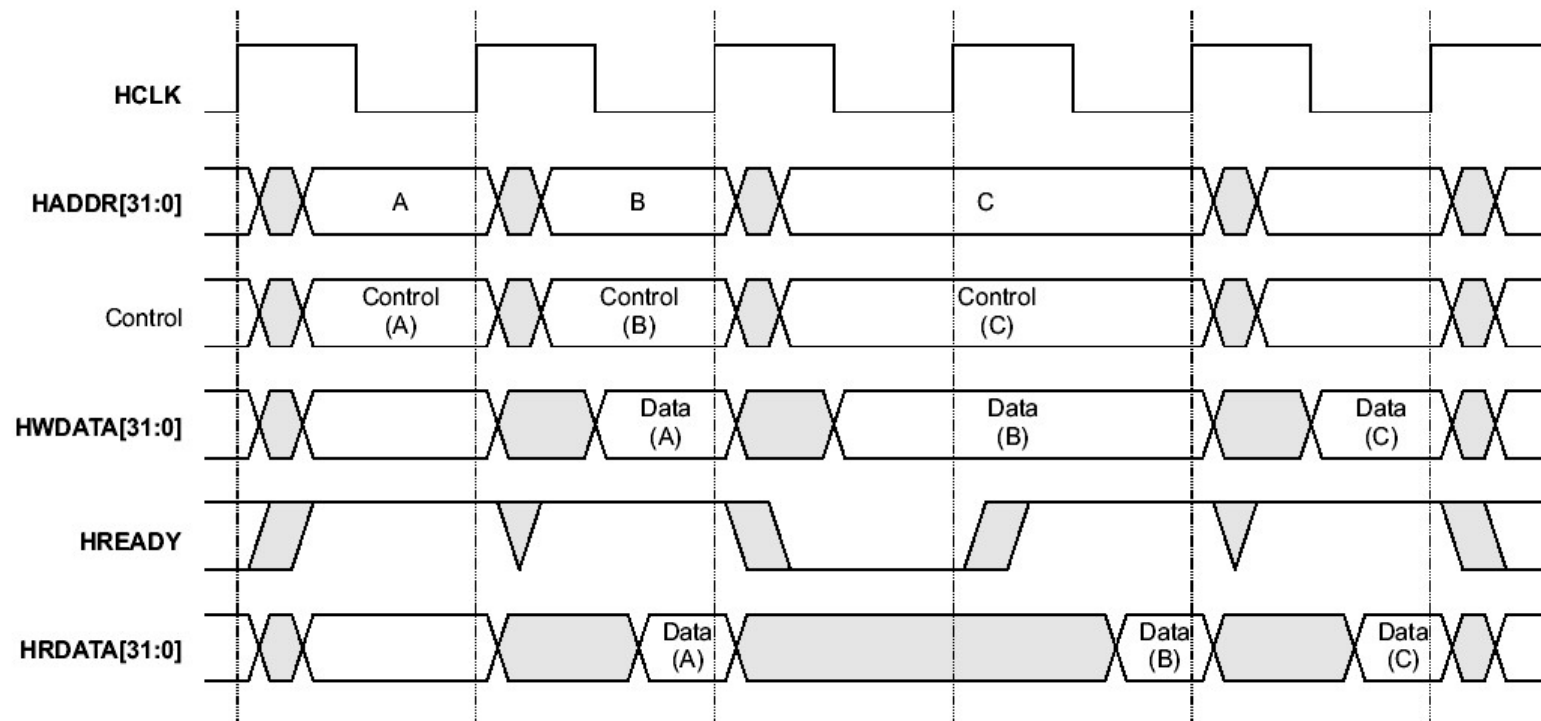
AHB Transfer with Wait State



† Figure extracted from: ARM Limited, *AMBA Specification Rev 2.0*, ARM IHI 0011A, May 13, 1999, figure 3-4, page 41.

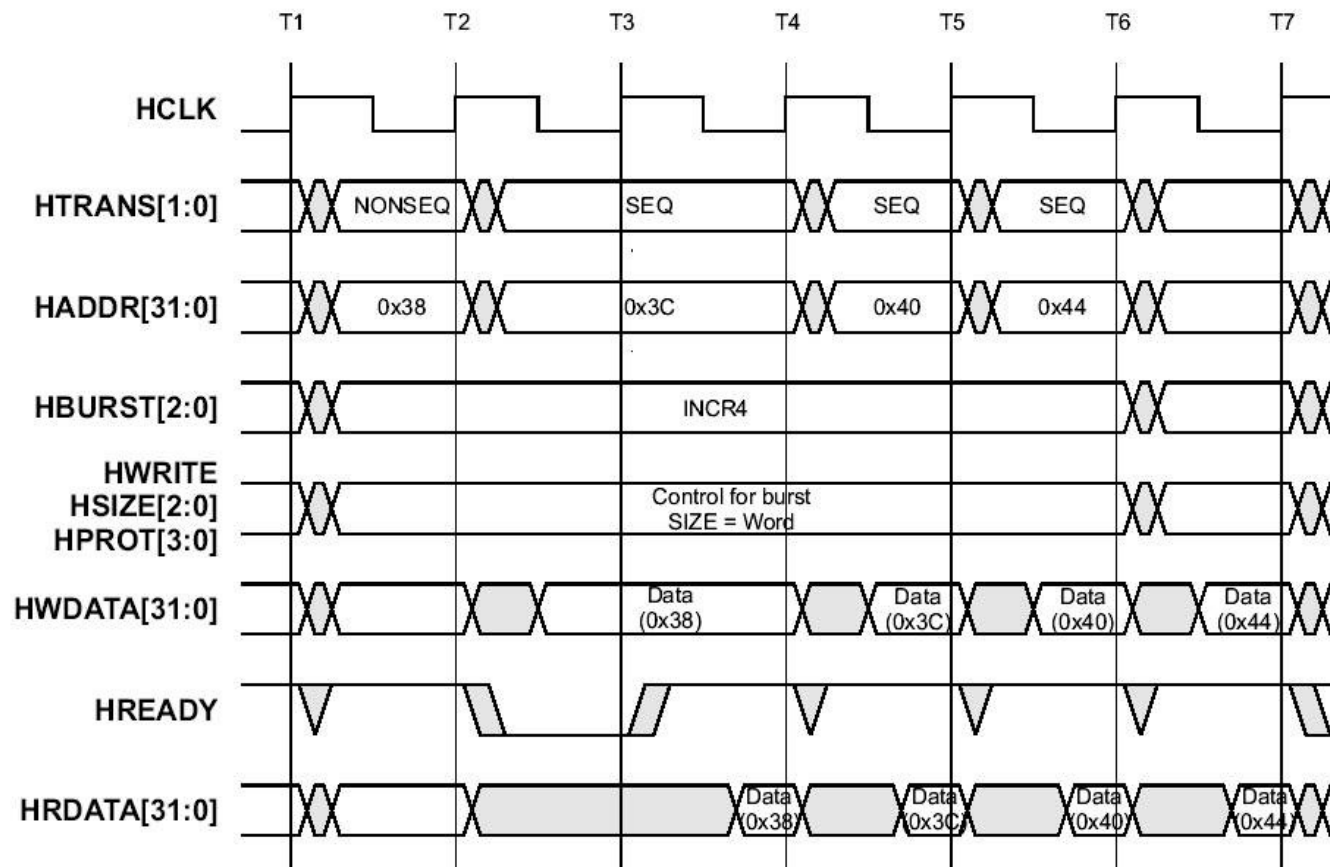
AHB Multiple Transfer

2-stage pipelined bus



† Figure extracted from: ARM Limited, *AMBA Specification Rev 2.0*, ARM IHI 0011A, May 13, 1999, figure 3-5, page 42.


Burst Transfer Example



† Figure extracted from: ARM Limited, *AMBA Specification Rev 2.0*, ARM IHI 0011A, May 13, 1999, figure 3-8, page 48.

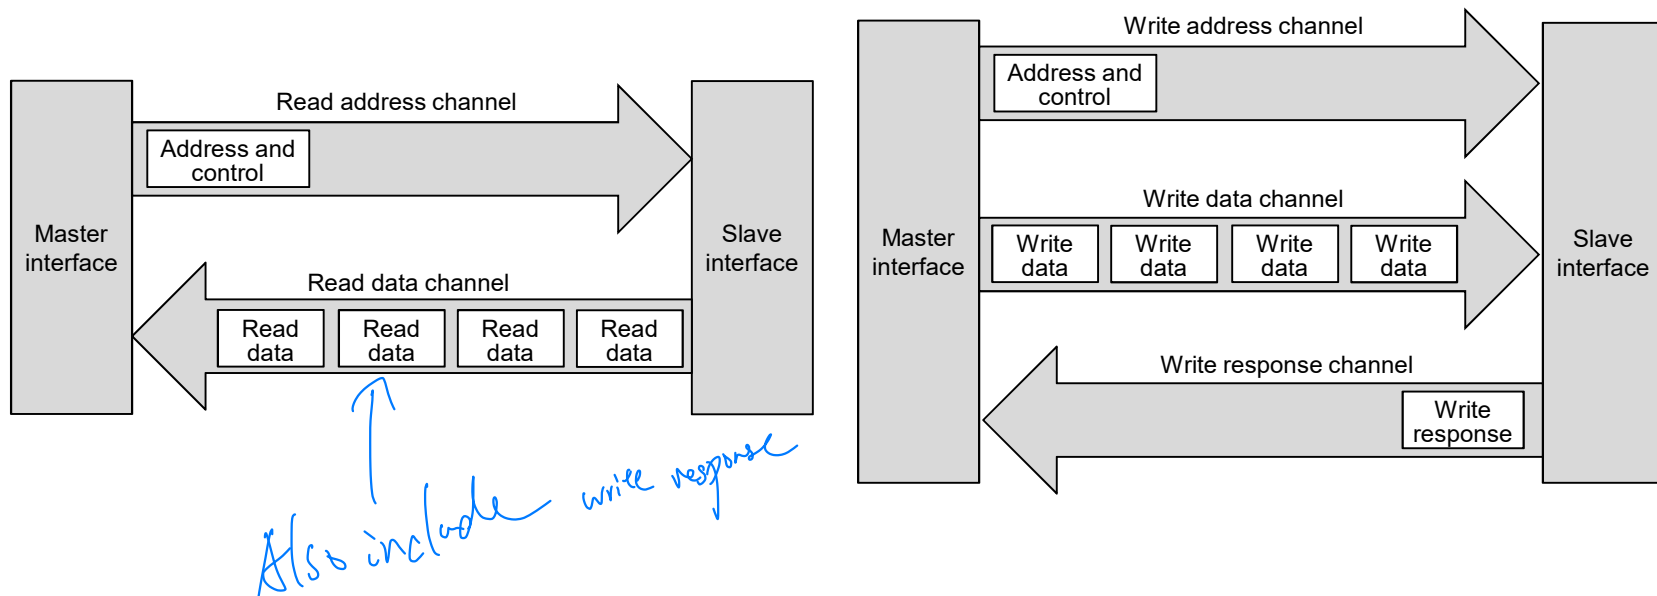
AXI4 Bus

- ❑ There are three types of AXI4 interfaces
 - AXI4 – for high-performance memory-mapped requirements
 - AXI4-Lite – for simple, low-throughput memory-mapped communication (e.g. for control registers)
 - AXI4-Stream – for high-speed streaming data



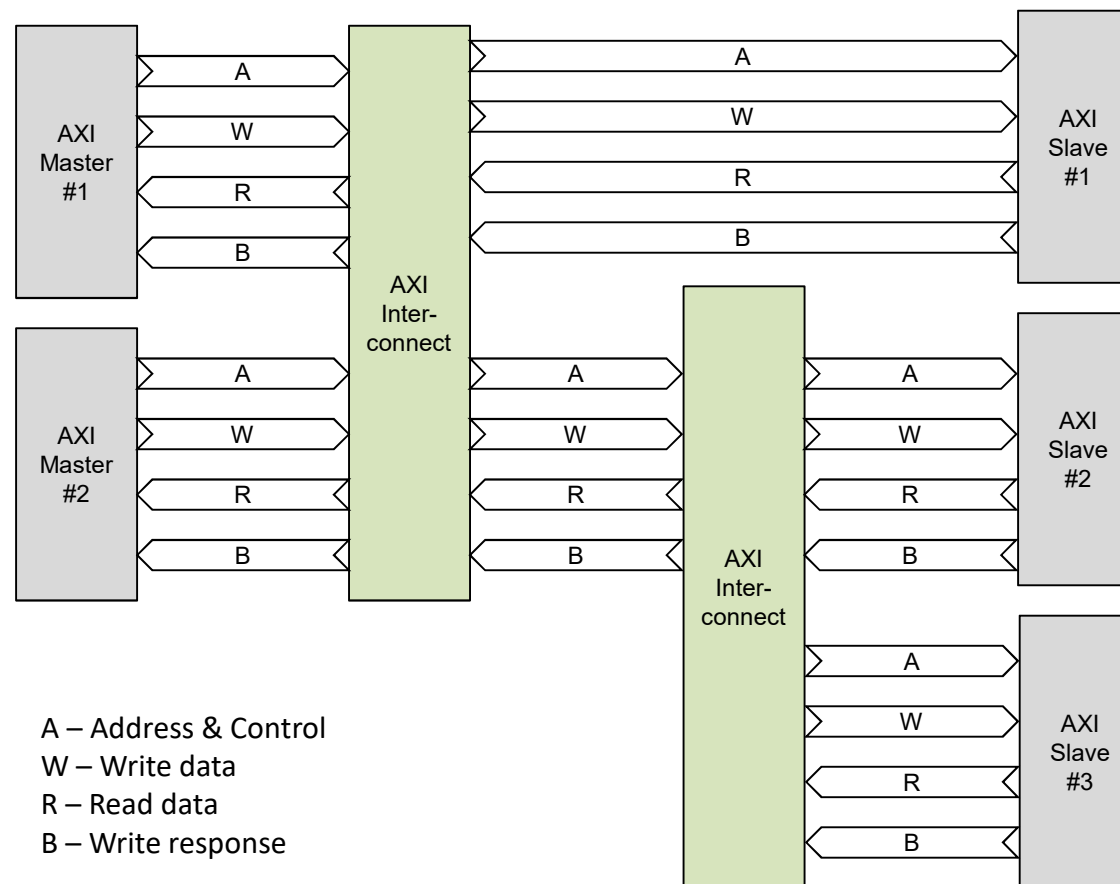
AXI4 and AXI4-Lite Interfaces

- ❑ Both AXI4 and AXI4-Lite interfaces consist of five uni-directional channels, read and write channels can operate simultaneously
 - Read Address Channel, Write Address Channel
 - Read Data Channel, Write Data Channel
 - Write Response Channel



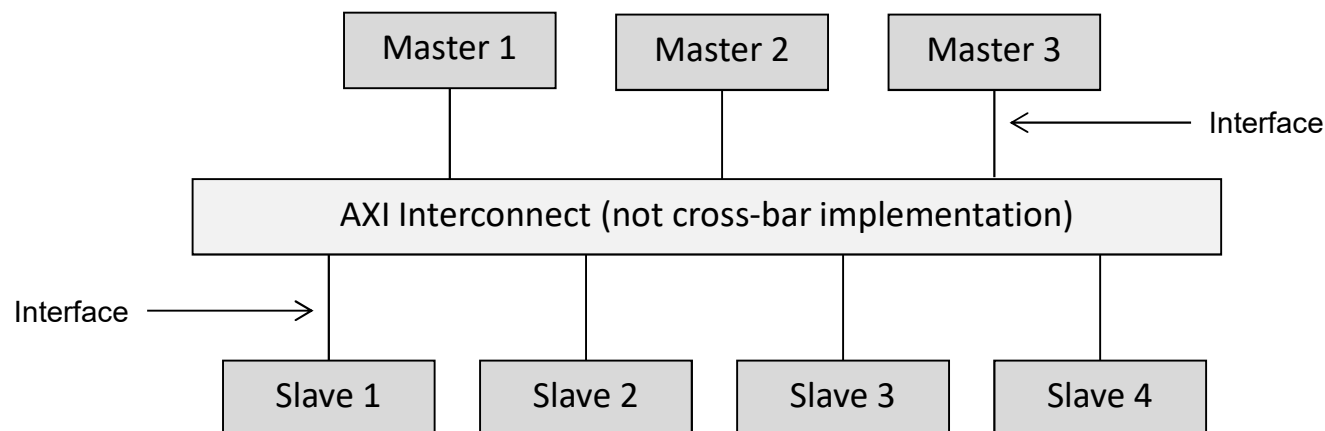
AXI Bus Topology Example (1/2)

- ❑ AXI devices do not share a common bus!



AXI Bus Topology Example (2/2)

- ❑ However, AXI can be used to construct a shared bus as well
 - The Wishbone bus protocol has similar concept
 - The interconnect implementation could be simplified to provide only one single transaction at any given time




AXI4 Advanced Features

- ❑ AXI4 allows a burst transaction of up to 256 data transfers; a burst transfer only needs to transfer a single starting address
- ❑ AXI4-Lite allows only one data transfer per transaction
- ❑ Features of AXI4
 - Data upsizing and downsizing
 - Multiple outstanding addresses
 - Out-of-order transaction processing
 - Local synchronization clock per master-slave pair
 - Insertion of register slices for improved timing closure

AHB vs. AXI Burst


□ AHB Burst

- Addresses and data are locked together
- HREADY controls intervals of address and data



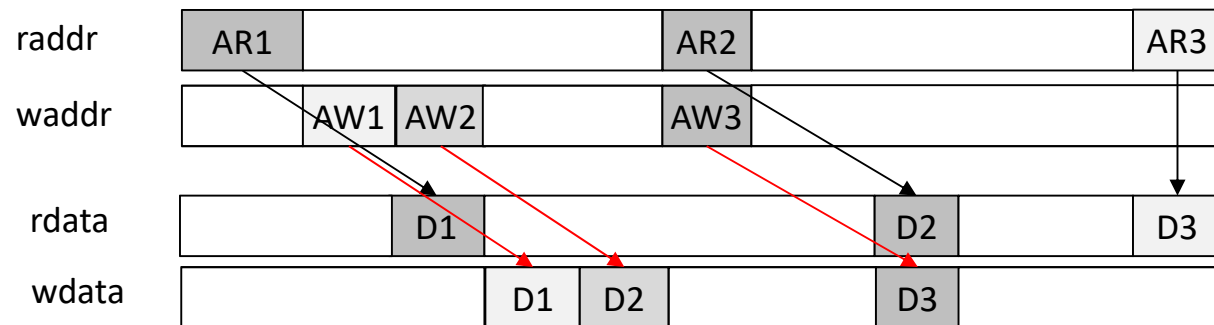
□ AXI Burst

- One address per burst



Interleaving Read/Write


- ❑ AXI burst allows simultaneous read/write transactions:



Out-of-Order Transmissions


❑ With AHB

- If one slave is very slow, all following transactions have to wait



❑ With AXI

- Fast slaves may return data ahead of slow slaves



Transaction ID's

- ❑ Out-of-order and interleaving transactions are possible because the AXI bus ties an ID signal to each address, data, and response signals.
- ❑ However, neither slaves and masters are required to honor the ID signals
 - Unless both the master and slave know about out-of-order transaction, you should not use this feature
 - A slave is mandatory to mirror the transaction ID from the master in BID (response ID) and RID (read data ID) of the return channels of the transaction

AXI4-Stream Features

- ❑ The AXI4-Stream bus is a single channel bus for transmission of unidirectional streaming data
 - Similar to the write data channel of the AXI4 bus, but there is no limit on the amount of data transfer per transaction
 - Used for point-to-point data transfer without a unified context of “address” between IP’s
 - More efficient than AXI4, but less flexible (e.g., random accesses are not allowed)


AXI Infrastructure IP's

- ❑ An infrastructure IP is used as a glue logic to construct a system, for example:
 - AXI interconnect IP
 - AXI FIFO's (for buffering/clock conversion)
 - AXI Direct Memory Access (DMA) engines
 - Register slices (for pipelining)

- ❑ These IPs are useful as “middlemen” in moving data around in a system, but not used as endpoints for data

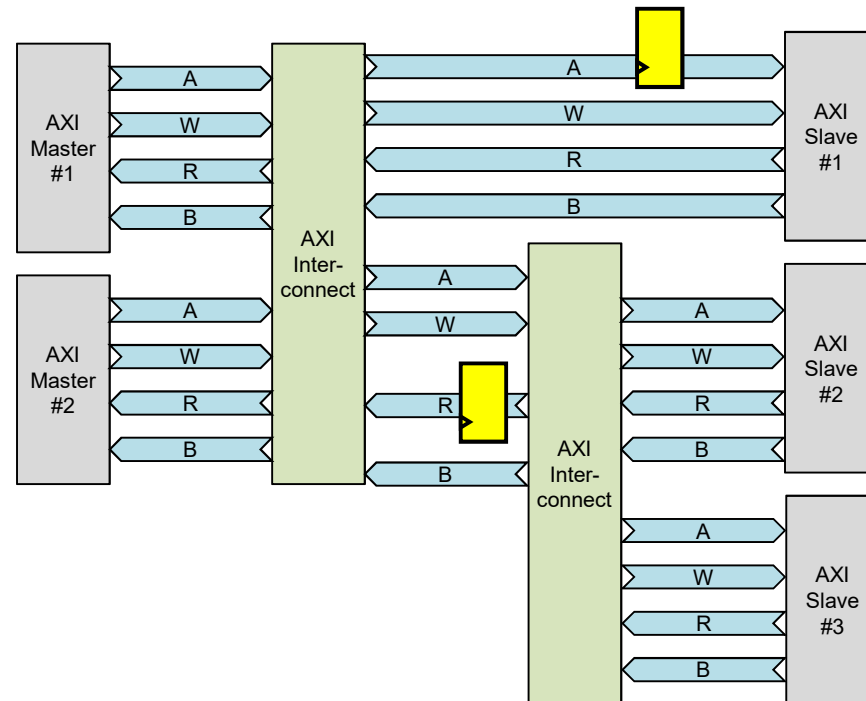
Example: AXI4 DMA Engine IP

- ❑ A DMA Engine IP is used to connect a AXI4-Stream interface (IF) to a AXI4 memory-mapped interface, and vice versa.



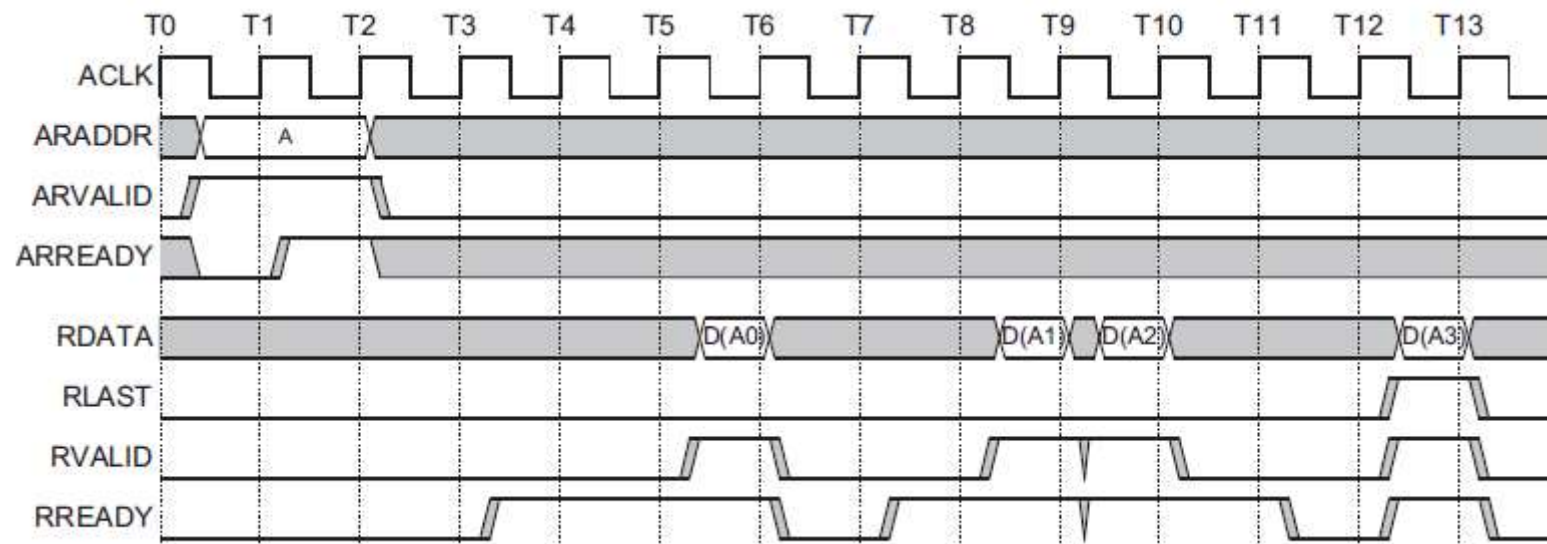
Example: Register Slices

- ❑ Each AXI channel transfers information in only one direction, and the architecture does not require any fixed synchronization
 - A register slice can be inserted at any point in any channel to increase working frequency, at the cost of additional latency



AXI Read Burst Example

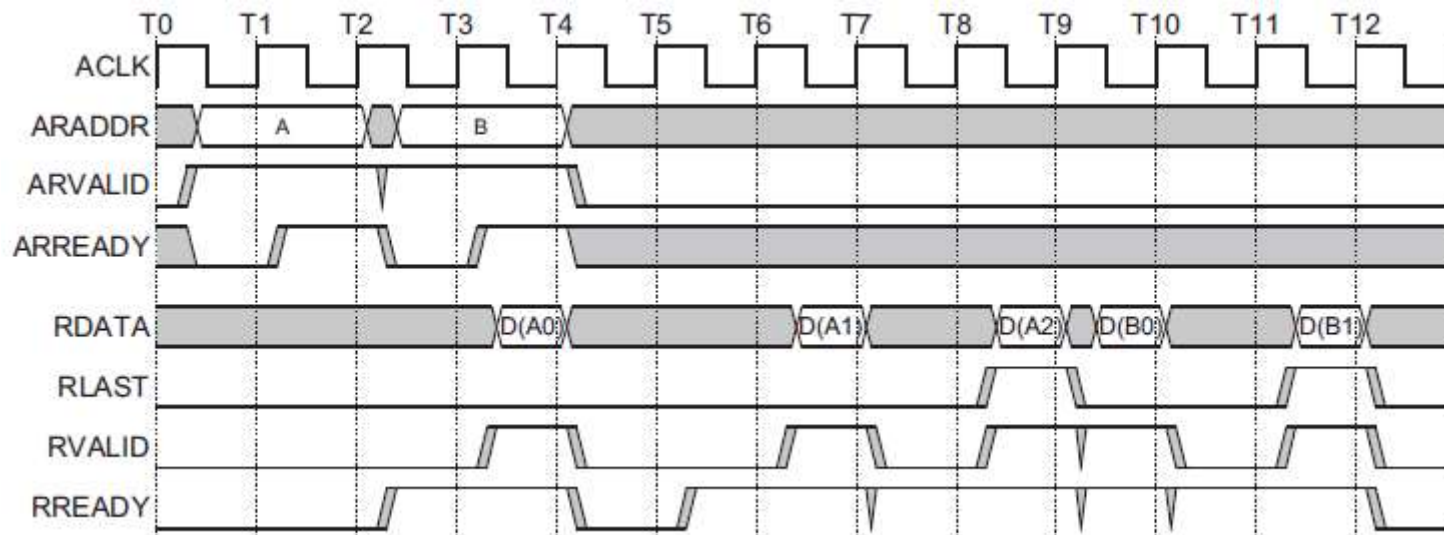
- A master performs a read burst of four data transfers:



† Figure extracted from: ARM Limited, *AMBA AXI Protocol v 1.0*, ARM IHI 0022B, March 19, 2004, figure 1-4, page 25.

Overlapping Burst Example

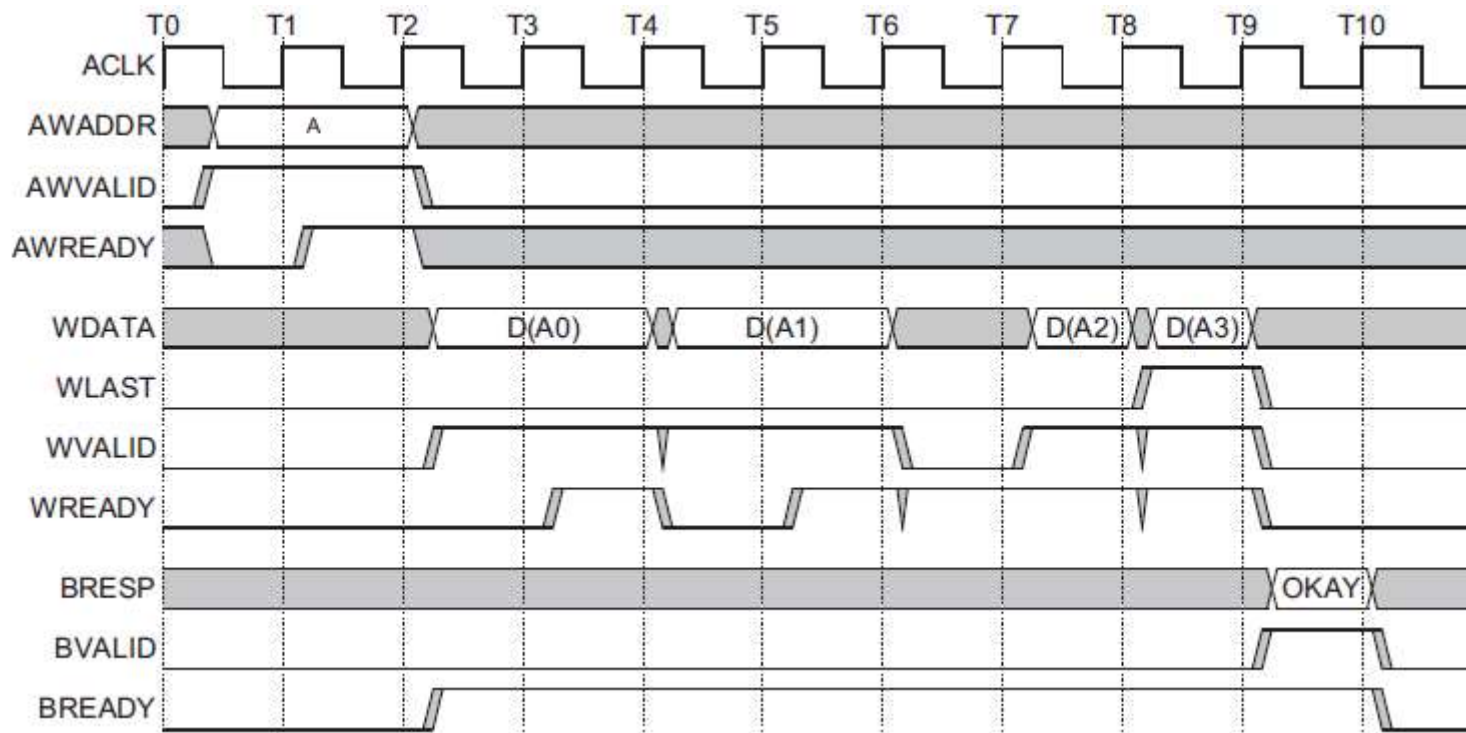
- ❑ A burst can start before the previous one is done:



† Figure extracted from: ARM Limited, *AMBA AXI Protocol v 1.0*, ARM IHI 0022B, March 19, 2004, figure 1-5, page 26.

AXI Write Burst Example

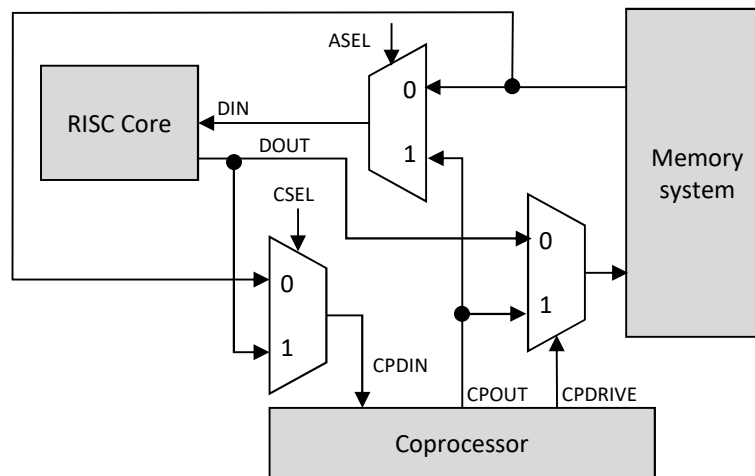
- A master perform a write burst of four data transfers:



† Figure extracted from: ARM Limited, *AMBA AXI Protocol v 1.0*, ARM IHI 0022B, March 19, 2004, figure 1-6, page 27.

Coprocessor Interface

- ❑ Traditional coprocessor interface:
 - ISA has coprocessor instructions (identify coprocessor ID)
 - Each coprocessor has its own set of registers
 - Coprocessors are not for high-bandwidth data processing
 - MMIO accelerators are more suitable for such tasks



Note:
Typical coprocessor does not control the external address bus – RISC CPU generates the required addresses.

Discussions

- ❑ Plug-and-play bus protocols enables reusable IP business
 - Small IP designers can sell their accelerator IPs to IC design houses as long as their IPs follows standard bus protocols such as the AXI bus
- ❑ On the other hand, integrating small circuit modules using AXI bus may be an overkill
 - Lightweight, proprietary bus protocols similar to AHB can be used internally within an IP