

HW#2 Branch Predictor Design



Chun-Jen Tsai
National Chiao Tung University
10/21/2022

Homework Goal

- ❑ This homework requires you to modify the Branch Prediction Unit (BPU) of Aquila to improve CoreMark
 - Analyze the current BPU first
 - Design a two-level predictor to improve the performance
- ❑ You have 2+1 weeks to implement this homework
 - You must upload your initial analysis report by 11/03, 17:00.
 - Students who have finished the analysis report get one extra week for preparing the final report. You should upload the final report and your code by 11/10, 17:00.

gshare — not very good --

Types of Branches

- ❑ There are three types of branches
 - Conditional forward jumps: for if-then-else statements
 - Conditional backward jumps: used in looping statements
 - Unconditional jumps: for function calls, or from bad coding
- ❑ The problem of branches:

```
318:  li    a0, 48
31c:  jal   ra, 102c
320:  addi   s1, s1, -1
324:  bne   s1, s3, 318
328:  li     s1, -1
32c:  srli   a5, s0, 0x1f
330:  add    s3, s0, a5
```

The next PC should be 0x102c, but
the Fetch unit does not know that until
two clocks later (after the Execute stage)

The Fetch unit does not even know what the
next PC should be until two clocks later (after
the Execute stage)

Aquila: Enable-branch-predictor

Dependencies of Fetch on Execute

- ❑ In a 5-stage pipeline, a branch instruction, after Fetch, may take up to two cycles to determine the next PC
 - The Execute is responsible for calculating the branch condition and update the PC
 - Do we have to stall the Fetch stage for the next instruction by two cycles?
- ❑ A branch predictor predicts the PC before Execute computes the condition expression
 - If the prediction is wrong, the pipeline has to be flushed before the Memory stage!

Static Branch Prediction

- ❑ Static branch prediction always make the same decision (forward/backward × taken/not taken)
- ❑ Implementation can be done by one of three methods
 - Hardwired into the processor pipeline
 - Assuming branch always taken, the Fetch must do a quick decode of the target PC
 - Assuming branch always not taken, then the $PC \leftarrow PC + 4$ (IBM).
 - Compilers generate the hint bit if the ISA supports it → PowerPC ISA.
 - Cooperation between the processor and the compiler, by following some register usage convention. For example,
 - “bne s1, s3, 318” suggests taken
 - “bne s1, s4, 318” suggests not taken

↙
e.g 2k reg : jump
e.g 2k+1 reg : nojump

Dynamic Branch Prediction

*transmeta :
modify the code while the program executes*

- ❑ The processor collects statistics **at runtime** of whether every branch instructions are taken or not
- ❑ The fetch unit fetches the **predicted** next instruction
- ❑ In the case of a misprediction, the pipeline has to be flushed to re-fetch the correct instruction
 - The penalty is high for a misprediction
 - The CPU states has not been changed upon misprediction

Stage #Cycles	Fetch	Decode	Execute	Memory	Writeback
1	BR	?	?	?	?
2	ADDI	BR	?	?	?
3	SLL	ADDI	BR	?	?
4	?	NOP	NOP	NOP	?

→ Next PC determined!

CPU state: register file


Branch Prediction Schemes

❑ One-level Predictor


- Uses a Branch History Table (BHT) indexed by the recent branch addresses
- When the fetch unit reaches a branch location, it gets the PC for the next instruction to fetch based on the BHT

❑ Two-level Adaptive Branch Prediction


- MCFarling's Two-Level Prediction (gshare, 1993).
- Call-stack predictor for function call returns



One-level Branch Predictor (1/2)




- ❑ One-level predictor is based on bimodal assumption:
 - The branch at a specific PC is either taken or non-taken most of the time
- ❑ For one-level branch predictor, we must determine:
 - How many address bits are used to index the BHT
 - How many bits are used to record the branch statistics
 - How many branch instructions are recorded in the BHT



One-level Branch Predictor (2/2)

- ❑ Aquila implements the simple 2-bit predictor
 - For each branch instruction, we record its branch likelihood with one of four possible states:




- The state changes after the execute stage determines whether the branch is taken or not.

BHT Implementation


** Aquila implements a fully associative (expensive)*

- ❑ The BHT is an associative memory
 - Data are retrieved using a TAG, instead of address
 - For BHT, TAG is the PC of the current branch instruction
- ❑ Theoretically, the BHT size should be large enough to accommodate all branch instructions in the program
- ❑ In Aquila, BHT size is specified by the parameter `ENTRY_NUM` in `bpu.v`
 - However, to actually change its size, you have to modify part of the Verilog code as well as the parameter.


One-way Set associative (direct mapped)



Two-way set associative

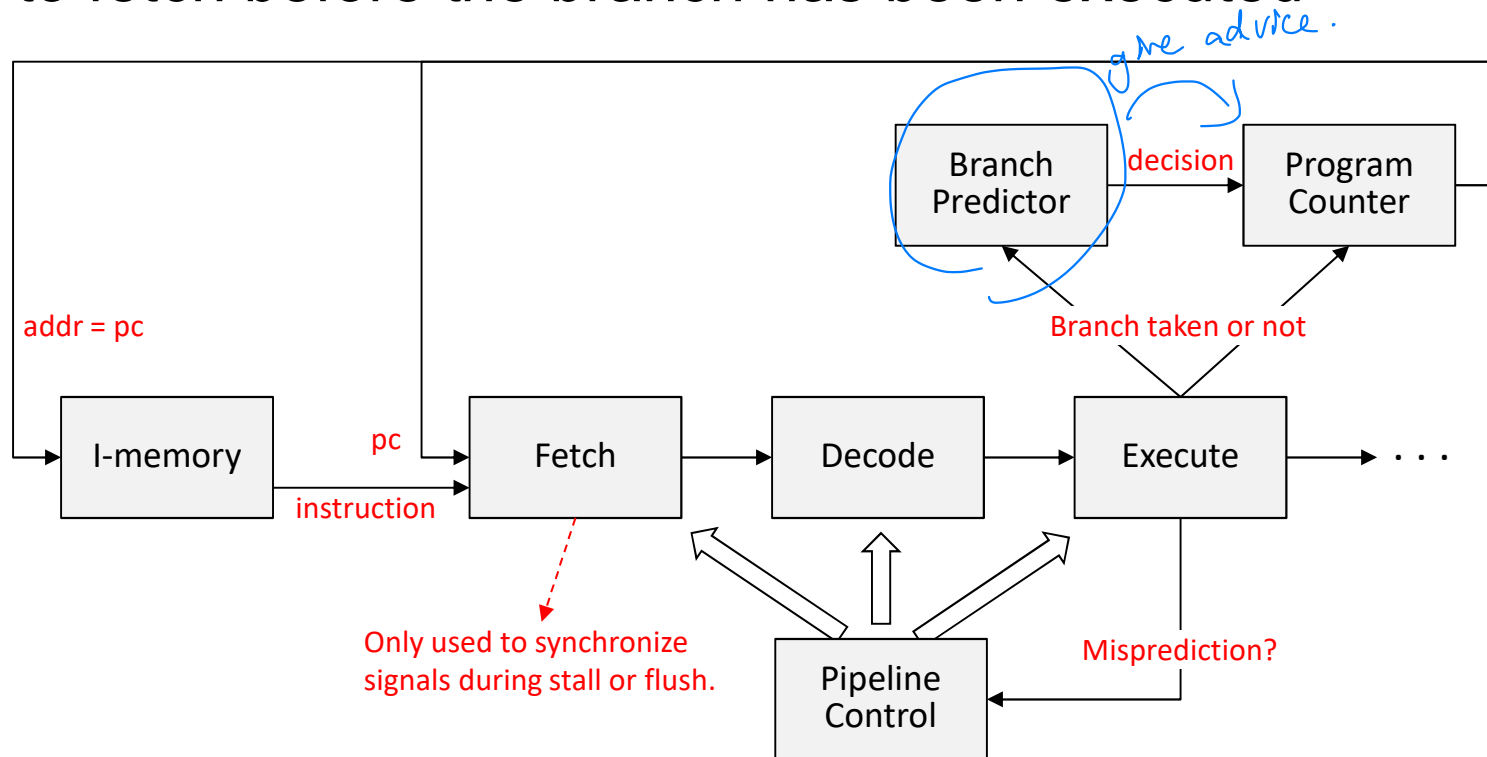


Fully associative

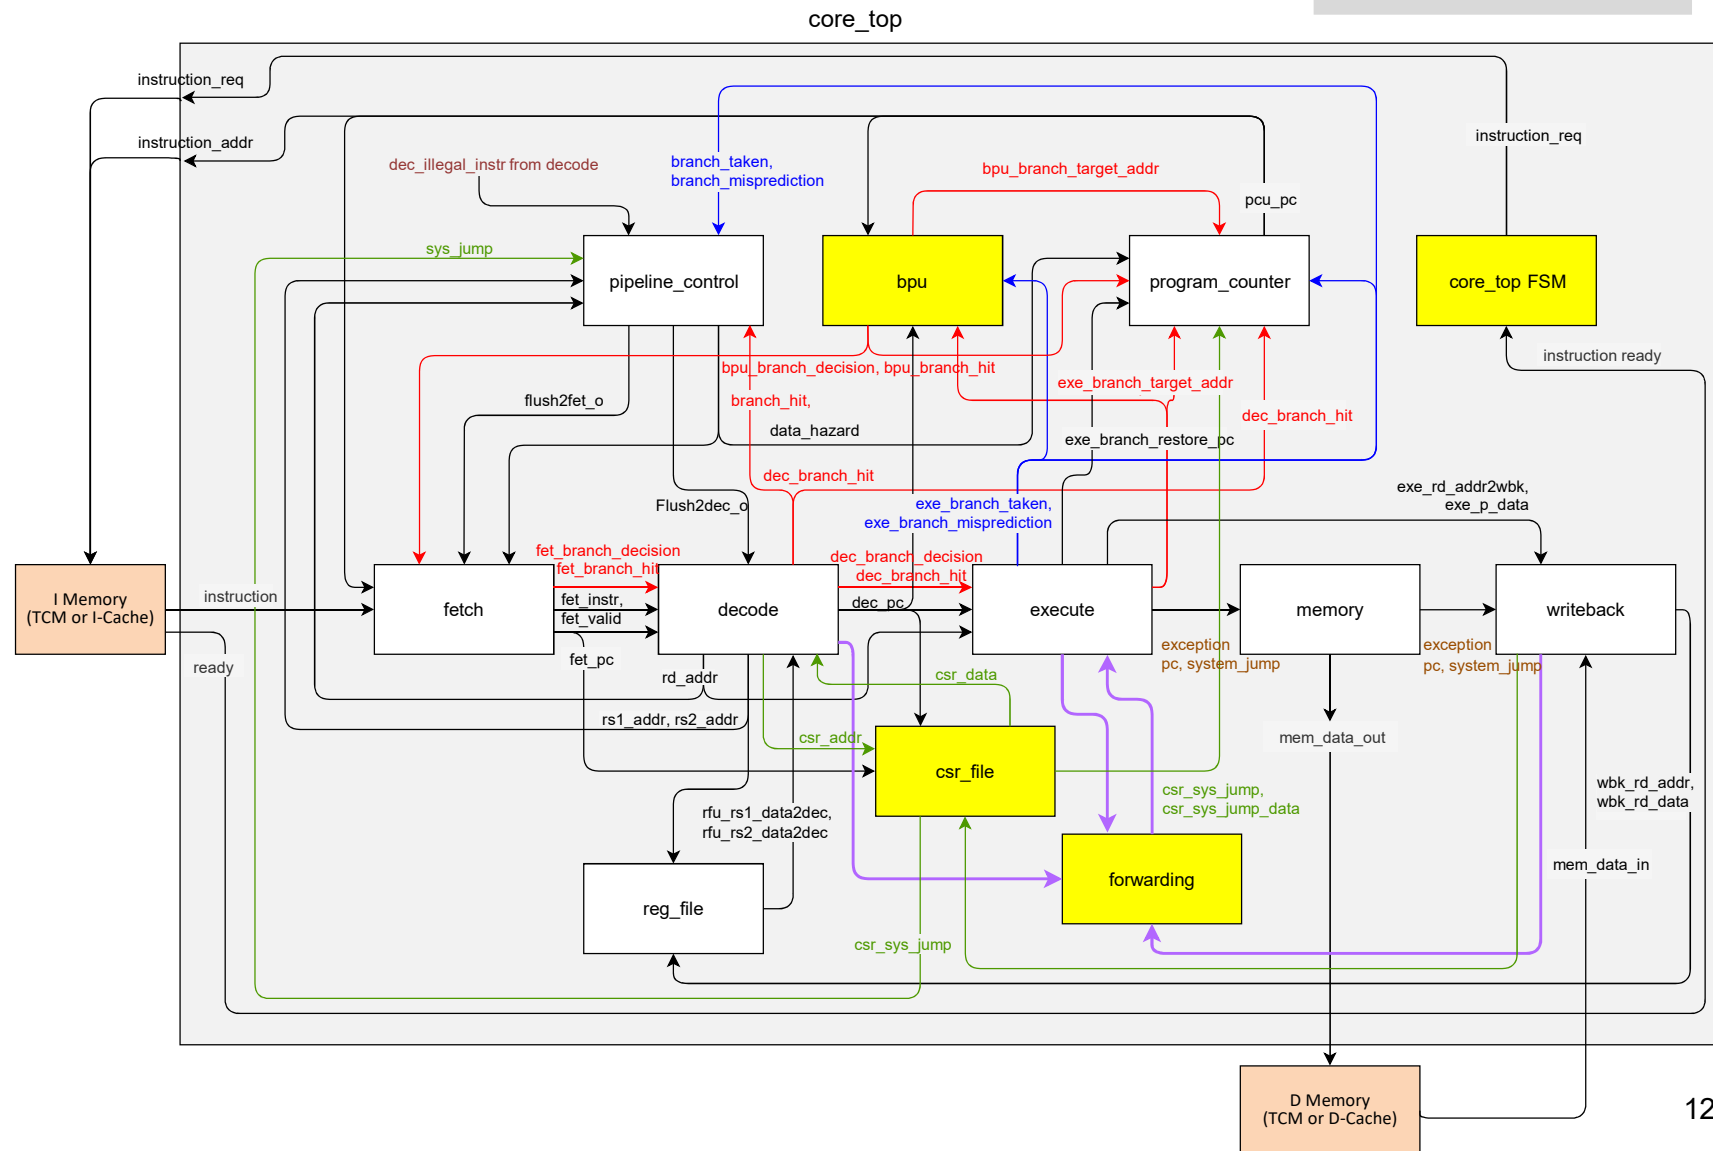


Branch Prediction Flow in Aquila

- ❑ A branch predictor tells the fetch unit which instruction to fetch before the branch has been executed

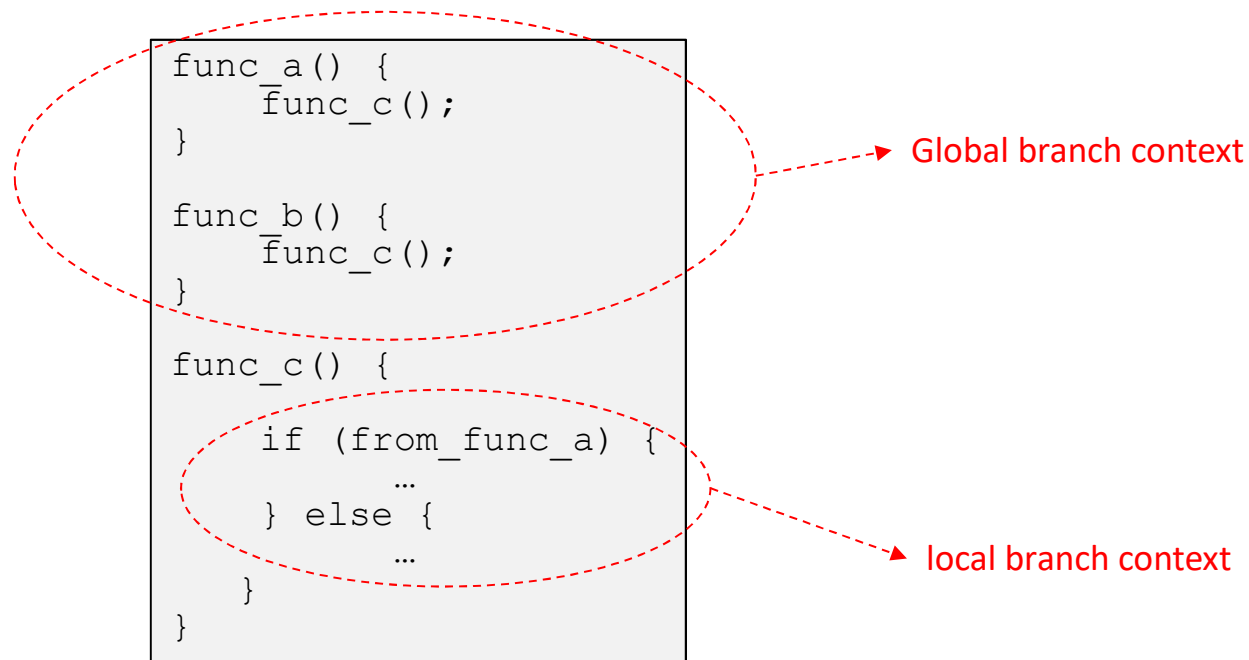


Aquila BPU Related Signals



Basic Ideas on Two-Level BPU

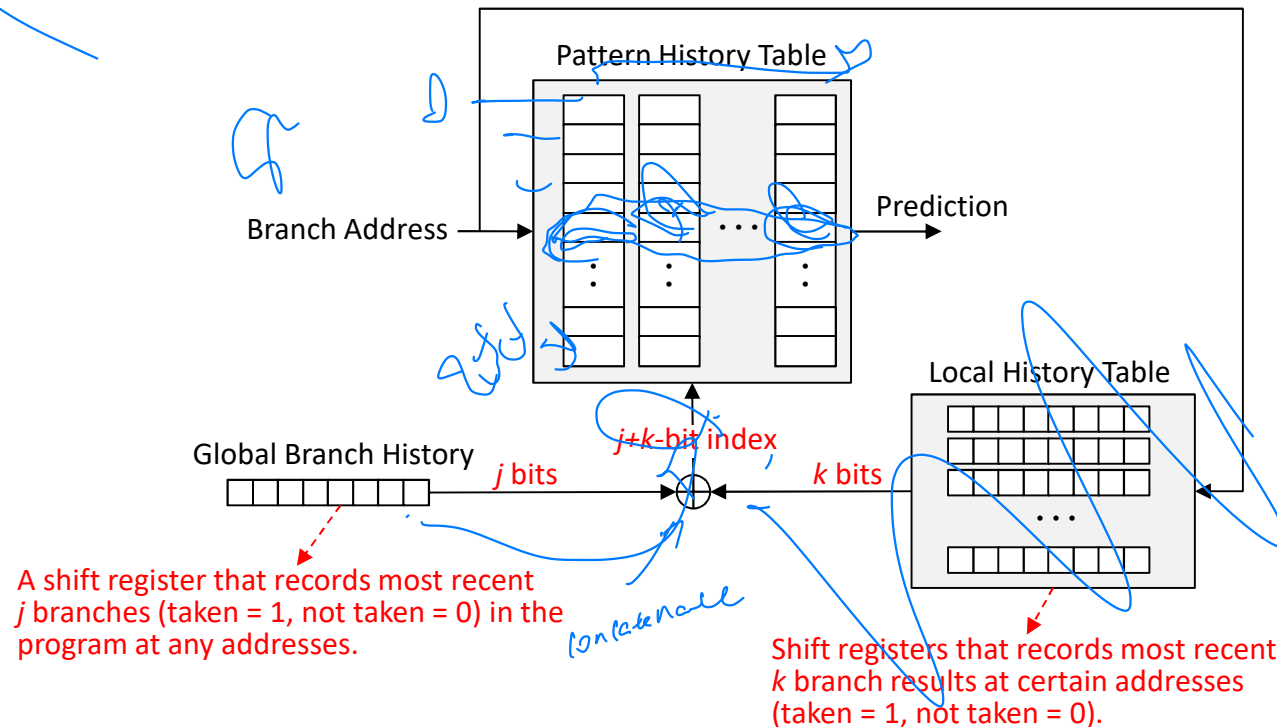
- ❑ One-level predictors lack global context information





- ❑ 2-level predictors consider both global & local contexts

Two-level Branch Predictor (1/3)




- ❑ A branch depends on the branch address as well as:
 - Nearby branches – recorded using Global Branch History
 - Longer history of the same branch – Local Branch History



PC




branch-label-hood



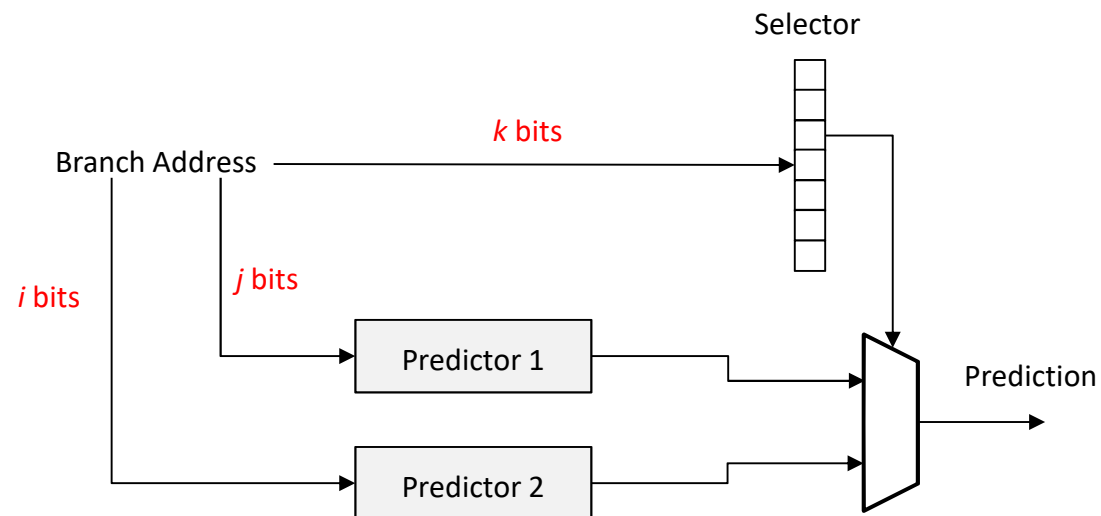
Two-level Branch Predictor (2/3)

- ❑ gshare: XOR address & history into the index bits
 - k-bit for global history would be a waste of registers
 - Global history and branch address can share the index bits using XOR operations



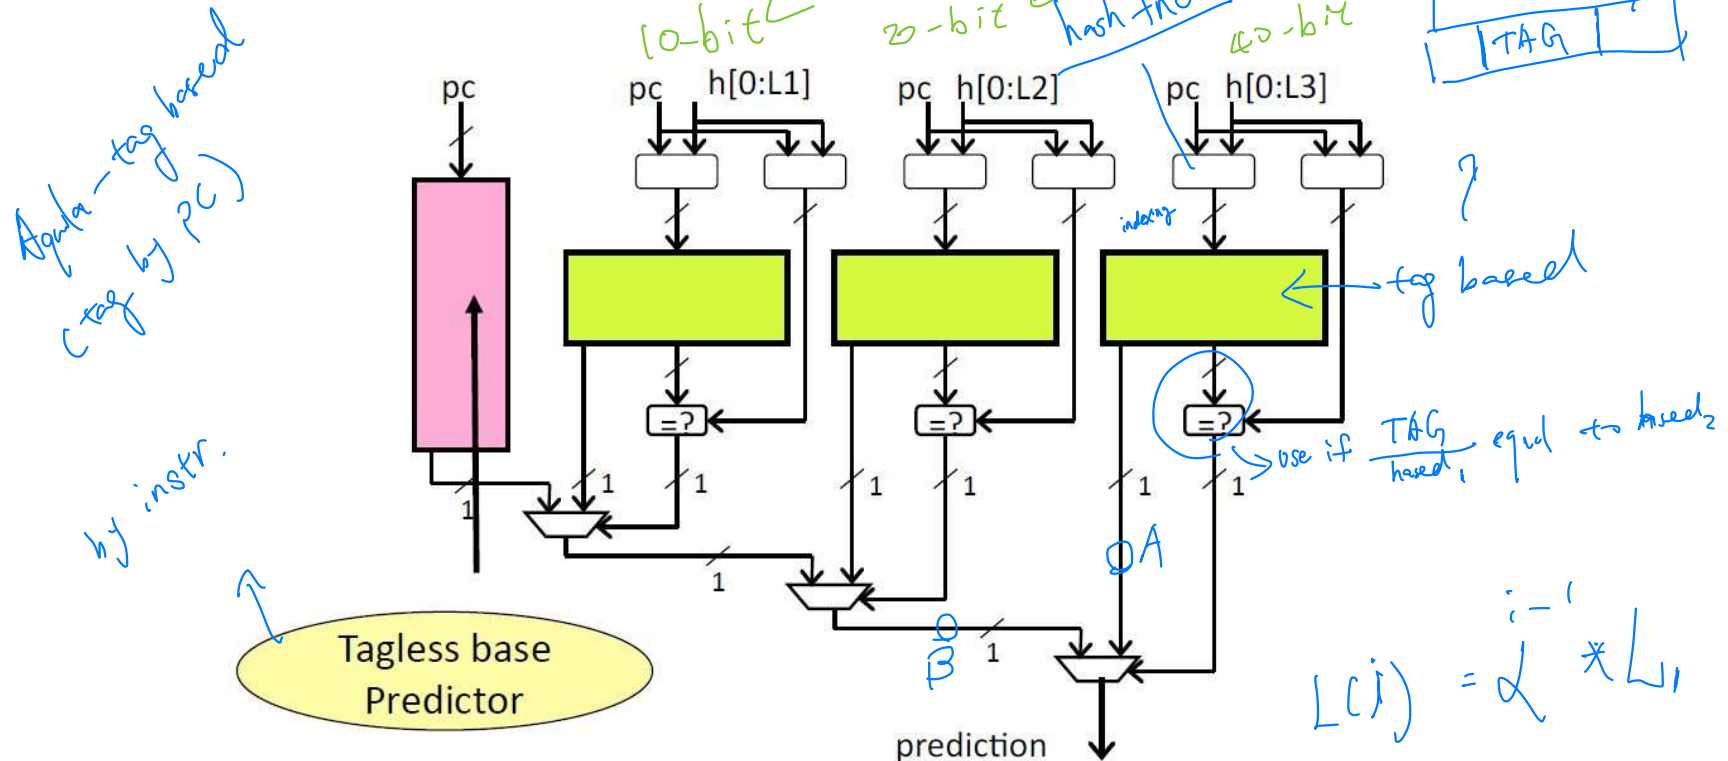
Two-level Branch Predictor (3/3)

- ❑ Different predictors work for different code patterns
 - Multiple predictors can be used to adapt to different code sections in the program:



TAGE Branch Predictors

- The state-of-the-art branch predictor is TAGE:



André Seznec, "A New Case for the TAGE Branch Predictor," MICRO 2011 : The 44th Annual IEEE/ACM Int. Symp. on Microarchitecture, Dec. 2011.

Your Homework (1/2)

- ❑ Part 1: Analysis report (60% of the credit):
 - Study the branch predictor in Aquila
 - Disable BPU or reduce BHT size see what happens to CoreMark
 - Analyze the branch statistics (e.g. hit rate and miss rate for different types of branches) of CoreMark
 - **Based on your analysis**, discuss your plan on how to use a two-level predictor to improve the performance
 - You report on this part should be no more than two pages.

Your Homework (2/2)

- ❑ Part 2: Two-level predictor (40% of the credit)
 - Try out some ideas on two-level predictors
 - **Extend your analysis report** to include new results and new discussions. The overall report size should be no more than three pages.